

# AgentGuard

## Phase 3 — Monetization & Scale

March 2026 · morphi101/agentguard · PRIVATE

### Where We Are

Phases 0, 1, and 2 are complete and tested. AgentGuard works. Here's what's been built:

| Phase   | What was built                                                                | Status    |
|---------|-------------------------------------------------------------------------------|-----------|
| Phase 0 | Shell interceptor, 30+ detection rules, log-based mode, 63 unit tests         | ✓<br>DONE |
| Phase 1 | Real PTY wrapper (node-pty), diff preview, config file, session summary       | ✓<br>DONE |
| Phase 2 | Local web dashboard, Telegram alerts (multi-recipient), context-aware scoring | ✓<br>DONE |
| Phase 3 | Cloud sync, team plans, policy templates, GitHub Action, billing              | →<br>NEXT |

### Phase 3 — What It Is

Everything built so far is local — it runs on your machine, stores data locally, and you're the only user. Phase 3 turns AgentGuard from a powerful personal tool into a **product you can charge for**.

It has five components:

## 1. AgentGuard Cloud

Today the audit log, rules, and config live only on your machine. Cloud adds a hosted backend so:

- Session history syncs across all your machines (Mac mini, laptop, servers)
- Custom rules defined on one machine appear everywhere automatically
- The dashboard is accessible from any browser, not just localhost:7429
- Team members see shared session history (with permission controls)

**Tech stack:** Node.js API (Express or Hono), PostgreSQL for session/event storage, Redis for real-time updates, deployed on Railway or Fly.io. Auth via magic link or GitHub OAuth.

## 2. Team Plan — Multiple Approvers

Today only you can approve or deny. With teams:

- Invite teammates by email — they get alerts on their own phones
- Define roles: who can approve **WARN** , who needs to see **CRITICAL**
- First-responder model: whoever responds first (approve/deny) wins
- Audit trail shows who approved what and when

This is the feature that makes AgentGuard a **team safety tool**, not just a solo developer tool. It unlocks the \$29/mo team tier.

## 3. Policy Templates

Instead of configuring rules from scratch, users pick a template that matches their context:

| Template           | Who it's for                     | Default posture                                |
|--------------------|----------------------------------|------------------------------------------------|
| <b>startup</b>     | Solo founders moving fast        | Permissive — only blocks CRITICAL              |
| <b>fintech</b>     | Financial apps, compliance-heavy | Strict — blocks WARN and above                 |
| <b>open-source</b> | Public repos with contributors   | Moderate — logs everything, blocks CRITICAL    |
| <b>ci-cd</b>       | Automated pipelines              | Auto-deny all risky ops, no interactive prompt |

## 4. GitHub Action

AgentGuard as a CI/CD step — audits the commands that run in your workflows and blocks dangerous ones before they reach production:

```
- name: AgentGuard scan
  uses: morphi101/agentguard@v1
  with:
    policy: fintech
    fail-on: CRITICAL
```

This opens the enterprise market — security teams that want a paper trail of every AI-generated command that ran in their pipelines.

## 5. Pricing Model

| Plan       | Price         | Machines  | Key features                                                   |
|------------|---------------|-----------|----------------------------------------------------------------|
| Free       | \$0 / forever | 1         | Local only, basic rules, CLI dashboard                         |
| Pro        | \$9 / month   | 3         | Cloud sync, all rules, Telegram alerts, policy templates       |
| Team       | \$29 / month  | Unlimited | Up to 10 users, roles, shared audit log, team approvals        |
| Enterprise | Custom        | Unlimited | SSO, SLA, GitHub Action, dedicated support, compliance reports |

At 100 Pro subscribers: **\$900 MRR**. At 100 Team subscribers: **\$2,900 MRR**. Break-even with infrastructure costs is around 30–40 subscribers.

## What Phase 3 Actually Requires

This is where the plan gets harder. Phase 3 is not just more code — it requires building and operating real infrastructure:

- **Backend server** — needs to run 24/7, handle auth, store data securely
- **Database** — PostgreSQL, backups, migrations, scaling
- **Auth system** — user accounts, sessions, team invites, permissions

- **Billing** — Stripe integration, webhooks, subscription management, failed payments
- **Domain + SSL** — agentguard.dev (check availability), HTTPS, CDN
- **Support** — users will have questions, bugs, and complaints
- **Documentation** — proper docs site, not just a README

Estimated build time for a functional Phase 3 MVP: **3–4 weeks of focused work**.

## My Honest Opinion

---

### ⚠ Don't build Phase 3 yet.

Not because it's a bad idea — it's a great idea. But because there are two things that need to happen first, and skipping them is how good products die quietly.

### First: Use it yourself for one real week

We built AgentGuard in one day. We tested it with a fake agent script. That's not the same as running `agentguard codex` on a real project, hitting real edge cases, and actually feeling whether it makes your life better or gets in the way.

If after one week of real use you find yourself turning it off — something is wrong with the product. If after one week you can't imagine working without it — you have something worth scaling.

**Action:** Use `agentguard codex` on a real project this week. Log every friction point.

### Second: Validate with the waitlist before building

The landing page at [bit.ly/4dFBaBY](https://bit.ly/4dFBaBY) is live. It's capturing signups. Before spending 3–4 weeks building cloud infrastructure, talk to 5–10 of those people.

Ask them:

- What AI agent do you use most?
- Has it ever done something destructive?
- Would you pay \$9/month for a tool that prevents that?
- Would you trust a cloud service with your session data, or do you need it to stay local?

That last question matters a lot. Some developers — especially those working on sensitive codebases — may never want their data leaving their machine. If that's the majority of the

market, the Free/local tier is where you stay, and Phase 3 looks very different.

 **Recommended sequence:**

- 1. Use AgentGuard on a real Codex session this week
- 2. Re-enable the landing page X posts (carefully — no source code leaks)
- 3. When you hit 50+ signups, reach out to 10 people directly
- 4. Make Phase 3 decision based on what you hear
- 5. If the signal is strong → build Phase 3
- 6. If people want local-only → focus on polish and npm distribution

**The risk of building too fast**

AgentGuard already works and solves a real problem. The biggest risk right now is not moving too slowly — it's building cloud infrastructure that nobody uses because you skipped the validation step. Three weeks of backend work at \$0 MRR is a high price to pay for a lesson that 5 user interviews would have taught you for free.

**If You Do Build Phase 3 — Timeline**

| Week   | Work                                                                                  |
|--------|---------------------------------------------------------------------------------------|
| Week 1 | Backend API + auth (GitHub OAuth + magic link), PostgreSQL schema, deploy to Railway  |
| Week 2 | Cloud sync in CLI (agentguard login, session upload), cloud dashboard, Stripe billing |
| Week 3 | Team plan (invites, roles, shared alerts), policy templates, npm publish v1.0         |
| Week 4 | GitHub Action, docs site, Product Hunt launch, HN "Show HN"                           |

**Your Notes**

---